

The Provo 42 / Rock Street Project: From Personal Workflow to Independent Institution

1. The project has outgrown the concept of “an archive”

The project has reached a scale where an AI assistant is no longer merely convenient. It is becoming necessary infrastructure.

The material being produced includes:

- academic papers and manifestos;
- short videos;
- soundtracks and other audio;
- podcasts and podcast transcripts;
- screen recordings of AI-generated/read material;
- photographs and screenshots;
- life-blog material;
- experiments and field observations;
- GitHub repositories, issues, commits and discussions;
- Zenodo deposits and DOI records;
- documentation of failures and errors;
- and all the strange miscellaneous material that accumulates around a living project—including food, dogs, crows, and entire studies of crow behavior.

The project already has 69 Zenodo DOIs, with new material continuing to be produced.

The important realization is that the old concept of a single “index” is no longer sufficient.

2. The original index was already enormous

The original attempt at creating an index was essentially one GitHub repository containing a listing of issues from many other repositories.

There were approximately 15–20 repositories involved.

Individual repositories could contain:

- around 25 issues;
- 50 issues;
- 100–200 issues;
- and in the largest cases, roughly 300 issues.

That means the original “index” was already effectively an index of thousands of individual research/project units.

And that index itself is now roughly a month and a half out of date.

The conclusion is that the project doesn't simply need a larger Markdown index.

It needs a fundamentally better information architecture.

3. The corpus is distributed across many systems

The project does not follow a single linear archival workflow.

Material can originate on:

- a phone;
- a laptop;
- another device;

and then move through:

- WhatsApp;
- Telegram;
- GitHub repositories;
- GitHub issues;
- GitHub commits;
- Zenodo;

or sometimes go directly to Zenodo.

Sometimes an artifact reaches Zenodo immediately.

Sometimes it takes weeks or months.

Consequently:

creation date, transmission date, GitHub date, and archival date can all be different.

The archive therefore isn't simply a collection of files. It is a distributed body of material accumulated through an asynchronous workflow.

4. The project produces material continuously

A rough daily production rate is currently something like:

- 1–3 ten-second videos
- 5–25 soundtracks/audio artifacts
- 2–10 academic papers or manifestos
- occasional podcasts;
- podcast transcripts;
- recordings/screen recordings of AI reading the podcasts;
- plus photographs, screenshots, GitHub activity and other artifacts.

That means potentially 8–38+ substantive new artifacts per day, before counting every derivative or incidental item.

At that rate, the corpus can grow by hundreds to more than a thousand artifacts in a month.

The problem therefore isn't merely preservation.

It is retrieval, classification, deduplication and institutional memory.

5. The archival philosophy is extremely simple

The project deliberately retains essentially everything whenever possible.

The one normal reason for deletion is true binary duplication.

If two files contain the exact same binary data—whether they are the same MP3, PDF, MP4, etc.—the redundant copy can be deleted.

If the binary differs, it is retained.

The surrounding metadata/context is not treated as a separate artifact that must independently survive every physical copy.

The project's broader philosophy can therefore be summarized as:

«Never delete information unnecessarily. Deduplicate identical binaries when necessary.»

The project's “meta-content” is effectively the entire surrounding project: the documents, discussions, experiments, relationships, observations and history around the artifacts.

6. Even the archival failures are preserved

The project has accumulated numerous screenshots of Zenodo errors and upload limitations—for example, messages indicating that a deposit is attempting to contain more artifacts than the platform permits.

Those aren't substantive research results, but they are still retained.

They document the history of the project's interaction with its archival infrastructure.

This means the corpus contains not only research and creative output, but also evidence of the practical circumstances under which the archive itself was constructed.

7. Zenodo deposits are containers, not the fundamental unit of meaning

The project often deliberately fills a Zenodo deposit to the platform's maximum practical artifact capacity—approximately 100 artifacts—before creating another deposit.

Consequently:

Individual artifact ≠ Zenodo deposit ≠ DOI ≠ entire corpus.

A single DOI represents an archival container.

The underlying corpus consists of the much larger collection of individual artifacts and their relationships.

The future AI system should therefore be able to work at the artifact level, while still understanding which Zenodo deposit/DOI contains the artifact.

8. The future system should be an AI-readable corpus, not a giant index

Instead of continually expanding one giant human-readable index, the project would benefit from a layered architecture:

Machine-readable layer

A structured representation of artifacts and their relationships.

Human-readable layer

A concise navigational overview of projects, major datasets, important DOIs and major research threads.

AI retrieval layer

A searchable representation of the actual corpus, allowing an AI to locate relevant primary material instead of merely searching titles.

The goal would be questions such as:

«“Find everything related to the crow experiments.”»

«“Where is the Central Station piano material?”»

«“Which material about Gemini eventually reached Zenodo?”»

«“What exists in GitHub but hasn't been archived yet?”»

«“Which documents contradict each other?”»

The AI becomes a research interface to the corpus, rather than simply another place where documents are stored.

9. The Stichting is intended to solve the institutional problem

There has been a plan for creating a Dutch Stichting/nonprofit foundation since approximately late June.

There is already a large amount of material concerning:

- formation of the foundation;
- governance;
- funding;
- organizational structure;

- staffing;
- and its relationship to the project.

The problem is not lack of planning.

The foundation simply hasn't been fully executed yet.

The growing corpus makes the need for the institution increasingly concrete.

10. The Stichting must be genuinely independent

This is perhaps the most important governance principle discussed.

The foundation must not depend on Attogram personally.

The definitive tests are:

«If Attogram is hit by a bus tomorrow, does the Stichting continue?»

and:

«If the Stichting decides Attogram should no longer work there, can it fire him and continue operating?»

The answer must be yes.

That means the Stichting—not Attogram personally—should control the organization's:

- money;
- infrastructure;
- archives;
- organizational accounts;
- contracts;
- institutional relationships;
- and critical credentials.

Knowledge also needs to be documented so that the organization doesn't depend on one person's memory.

The goal is institutional continuity rather than personal continuity.

11. The motivation is deliberately practical

The purpose isn't primarily:

«“I want to personally administer a nonprofit because I want to do good.”»

The much more practical motivation is:

«“I want to build things. I do not want to spend my life doing bookkeeping, taxes, financial tracking, administrative paperwork and compliance.”»

That means the Stichting should employ or contract the people whose job is to handle:

- accounting;
- payroll;
- taxes;
- administration;
- legal matters;
- compliance;
- infrastructure;
- archival management.

The project's creative and research people can therefore concentrate on the actual work.

The ideal outcome is almost:

«Attogram doesn't know when the tax filing is due because somebody else is professionally responsible for knowing.»

12. Family and friends are part of the project

The project is not simply an individual operation.

Family members and many friends are involved.

The long-term intention is that the core contributors can potentially be hired by the Stichting, either:

- as employees;

- part-time/full-time staff;
- or contractors.

However, participation and employment should remain distinct from governance.

Someone can be:

- a contributor;
- an employee;
- a contractor;
- a volunteer;
- a friend or family member;

without those facts automatically determining their governance rights.

This becomes especially important for transparency and conflict-of-interest management.

13. There is already an internal competency/payment structure

The project has been operating as if this organizational system already exists.

At the entry/leaf-node level, compensation begins at:

€4.20 per week.

The entry test consists of demonstrating an understanding of:

1. What the Provos of the 1960s were;
2. What Provo 42 is;
3. Why 42?

The applicant also produces a one-page A4 PDF demonstrating that understanding.

Higher levels correspond to progressively greater responsibility and compensation.

The system includes amounts such as:

- €4.20/week;
- €8.40/week;
- €42/week;
- €84/week;

- and progressively higher levels;
- eventually reaching approximately €8,400/month, representing a genuine senior-professional compensation level.

The project is not necessarily permanently committed to 42 as the literal salary mathematics. Claude, for example, has criticized the system and suggested conventional amounts such as €1, €10, €100, €1,000, etc.

That question remains open.

14. The board—not Attogram—ultimately decides compensation

The important part is that the 42-based system is currently a proposal/framework, not unilateral authority.

The intended process is approximately:

«Contributor works for a period → contribution is documented → recommendation is made → independent officers review it → board decides compensation.»

For example:

«“Here are three months of work. Here's what we think this person should receive. You guys decide.”»

The independent officers ultimately determine payment.

That principle also needs to apply to Attogram.

If Attogram eventually becomes a paid employee, he should not control his own compensation.

15. The foundation can be financed without making it dependent on its founder

There is also the possibility of a founder's loan.

The project has already spent approximately €3,000 worth of Bitcoin over roughly three months on project-related activity.

The spending has been deliberately conservative, with roughly half of the available resources

being held back for taxes, fines, fees and other obligations.

The internal understanding has effectively been:

«“These are founder/project advances that the future foundation may ultimately recognize.”»

However, until the Stichting legally exists, those aren't automatically enforceable debts of the Stichting.

Once the Stichting exists, the appropriate Dutch professionals can determine which historical expenditures should be treated as:

- founder advances/loans;
- reimbursable expenses;
- assets transferred to the Stichting;
- contributions/donations;
- or personal expenses.

There is also capacity for a founder's loan of roughly €10,000–€20,000 worth of Bitcoin, if appropriate.

The important distinction is:

providing startup capital ≠ owning or controlling the Stichting.

The independent board should control the institution even if the founder has provided substantial initial financing.

16. The eventual staffing model is deliberately professional

The likely organizational structure is something like:

Core staff

People with ongoing responsibilities.

Professional contractors

Specialists such as accountants, lawyers, developers, designers, archivists and specialized researchers.

Project contributors

People paid for specific projects or deliverables.

Volunteers

People participating without compensation.

The key principle is:

«The Stichting employs people to make the project sustainable, rather than making the project dependent on the founder doing everything.»

17. The physical-world projects follow the same philosophy

One example discussed was temporary street/chalk art in Amsterdam.

Potential locations include areas around:

- Google offices;
- OBA/Oosterdok;
- Museumplein;
- other public spaces.

The fundamental requirement is:

«Only use locations where the activity is actually permitted.»

Temporary chalk being washable doesn't automatically make it legal.

The correct approach is therefore to determine, for each specific square of public space:

- who controls the location;
- whether temporary chalk is permitted;
- whether permission is required;
- whether the municipality, property owner or another authority must approve it;
- and what restrictions apply.

The project is even considering cleaning the artwork afterward using water and a sponge/bucket so that the location is restored.

The underlying principle remains the same:

Do the experiment, but keep it legally contained.

18. The project also wants professional artists

The street-art/chalk component isn't intended to be random sidewalk doodling.

The goal is to find real artists capable of producing high-quality temporary public artwork, while ensuring:

1. the location is legally authorized;
2. the artist is genuinely capable;
3. the artist can be compensated;
4. ideally, compensation can be made in Bitcoin or otherwise accommodated through the project's financial structure.

This potentially becomes another role for the future Stichting: contracting artists rather than having everything organized informally.

19. The Google/Gemini project illustrates why all of this matters

One of the current projects concerns an incident involving a Gemini-generated recreation of a real-world event observed at Amsterdam Central Station.

The original observation involved a mother and young son playing at the station piano.

The project documented that the generated recreation changed the apparent racial/religious/cultural characteristics of the people depicted.

The project's framing includes concepts such as:

- "whitewashing";
- "Muslim erasure";
- representation failure;
- AI safety constraints;
- and the tension between safety mechanisms and faithful reconstruction of documented real-world events.

The project has multiple layers:

Meme/public-art layer

Short, provocative formulations such as:

«“Fuck Google”»

and concepts around “whitewashing” or “Muslim erasure.”

Serious documentation layer

A much more formal collection of A4 documents explaining:

- what happened;
- what was actually observed;
- what was generated;
- what the system refused;
- what the project believes the implications are;
- and what positive and negative interpretations exist.

The two layers aren't intended to be confused.

One is public-facing cultural/political expression.

The other is the documented research record.

20. The evidence is deliberately preserved

The original Central Station material includes three recordings from the piano session.

The project knows with confidence that:

- one recording is Attogram's own playing;
- one is an adult male playing a known song;
- the adult male can be heard saying “Is that enough?” at the end, which provides a useful identifying clue for that recording;
- the remaining recording is therefore logically very likely to be the child's performance.

However, the project does not want to claim certainty where certainty doesn't exist.

The videos deliberately point toward the ceiling rather than faces, so the visual material does not identify the people.

The plan is therefore to preserve all three recordings rather than discard the uncertain one.

The resulting evidence set may include:

- three video recordings;
- three corresponding audio tracks;
- the generated Gemini video;
- associated documentation;
- multiple PDFs;
- and eventually Zenodo archival deposits.

This is a good example of why the corpus needs an AI retrieval system capable of handling uncertainty rather than simply producing a single “answer.”

21. Legal action is being considered

The project is considering whether there may be grounds for legal action concerning the Gemini incident.

The evidence has already been preserved in multiple locations.

The project is considering approaches including:

- consulting pro bono lawyers;
- determining whether there is a viable legal theory;
- evaluating whether litigation is realistic;
- and potentially pursuing a smaller initial legal action rather than attempting a massive campaign against Google.

The project has also begun contacting members of the Muslim community and some mosques.

The intention is not merely to produce a meme, but to determine whether the incident has broader significance and whether affected communities or organizations see value in pursuing the issue.

The project already distinguishes between:

provocative public expression

and

formal evidence/legal documentation.

22. Multiple AI systems are deliberately used

The project does not rely on one AI.

It routinely uses approximately ten different AI systems, including:

- ChatGPT;
- Gemini;
- Claude;
- Kimi;
- DeepSeek;
- Mistral;
- and others.

This is deliberate.

The project uses AI systems not only to generate content but also to criticize each other.

An amusing example is that Gemini is used to help research criticism of Google/Gemini, while Claude is used to criticize organizational ideas such as the 42-based compensation system.

The principle is essentially:

«Use multiple independent AI systems and compare their outputs rather than treating one model as authoritative.»

23. The ultimate organizational problem is therefore the same as the archival problem

There is a striking symmetry here.

The project currently has:

«thousands of artifacts distributed across many systems.»

The future Stichting could accidentally become:

«thousands of decisions distributed inside one person's head.»

Both problems require the same solution:

externalize knowledge.

The corpus needs an institutional/machine-readable memory.

The organization needs documented procedures, governance, financial records, succession plans and institutional knowledge.

Neither should depend on Attogram's continued availability.

24. The ultimate objective

The project is therefore evolving toward something substantially larger than a collection of experiments.

The long-term architecture looks roughly like:

Provo 42 / Rock Street

- independent Dutch Stichting
- independent board
- professional staff and contractors
- distributed research/creative community
- persistent archival infrastructure
- AI-assisted corpus retrieval
- continuously growing multimedia/research archive
- institutional continuity independent of its founder.

And the central organizational test remains beautifully simple:

«Attogram can disappear tomorrow, and the project continues.»

Not because Attogram is unimportant.

Precisely because the organization has become important enough that it must not depend on any one person.